

Decision and Estimation in Data Processing

Chapter II. Elements of Signal Detection Theory

II.1 Introduction

Introduction

- ▶ Signal detection = the problem of deciding which signal is present from 2 or more possibilities
 - ▶ one possibility may be that there is no signal
- ▶ Based on **noisy** observations
 - ▶ signals are affected by noise

The model for signal detection




Figure 1: Signal detection model

► Contents:

- Information source: generates messages a_n with probabilities $p(a_n)$
- Modulator: transmits a signal $s_n(t)$ for message a_n
- Channel: adds random noise
- Sampler: takes samples from the signal $s_n(t)$
- Receiver: **decides** what message a_n has been transmitted

Practical scenarios

► Data transmission

- constant voltage levels (e.g. $s_n(t) = \text{constant}$)
- PSK modulation (Phase Shift Keying): $s_n(t) = \text{cosine}$ with same frequency but various initial phase
- FSK modulation (Frequency Shift Keying): $s_n(t) = \text{cosines}$ with different frequencies
- OFDM modulation (Orthogonal Frequency Division Multiplexing): particular case of FSK

► Radar

- a signal is emitted; if there is an obstacle, the signal gets reflected back
- the receiver waits for possible reflections of the signal and must decide
 - no reflection is present -> no object
 - reflected signal is present -> object detected

Generalizations

- ▶ Decide between more than two signals
- ▶ Number of observations:
 - ▶ use only one sample
 - ▶ use multiple samples
 - ▶ observe the whole continuous signal for some time T

II.2 Detection of constant signals

Detection of a constant signal, 1 sample

- ▶ Simplest case: detection of a constant signal contaminated with white normal noise, using 1 sample
 - ▶ two messages a_0 and a_1
 - ▶ messages are encoded as constant signals
 - ▶ for a_0 : send $s_0(t) = 0$
 - ▶ for a_1 : send $s_1(t) = A$
 - ▶ over the signals there is additive noise
 - ▶ receiver takes just 1 sample
 - ▶ decision: compare sample with a threshold

Threshold-based decision

- ▶ The value of the sample taken is $r = s + n$
 - ▶ s is the true underlying signal ($s_0 = 0$ or $s_1 = A$)
 - ▶ n is a sample of the noise
- ▶ n is a (continuous) random variable
- ▶ r is a random variable also
 - ▶ what distribution does r have compared to n ?
- ▶ Decision is taken by comparing with a threshold T :
 - ▶ if $r < T$, take decision D_0 : decide the true signal is s_0
 - ▶ if $r \geq T$, take decision D_1 : decide the true signal is s_1

Hypotheses

- ▶ Receiver chooses between **two hypotheses**:
 - ▶ H_0 : true signal is s_0 (a_0 has been transmitted)
 - ▶ H_1 : true signal is s_1 (a_1 has been transmitted)
- ▶ Possible results
 1. **Correct rejection**: no signal present, no signal detected.
 - ▶ Decision D_0 when hypothesis is H_0
 - ▶ Probability is $P_n = P(D_0 \cap H_0)$
 2. **False alarm**: no signal present, signal detected (error)
 - ▶ Decision D_1 when hypothesis is H_0
 - ▶ Probability is $P_{fa}P(D_1 \cap H_0)$
 3. **Miss**: signal present, no signal detected (error)
 - ▶ Decision D_0 when hypothesis is H_1
 - ▶ Probability is $P_m = P(D_0 \cap H_1)$
 4. **Hit**: signal present, signal detected
 - ▶ Decision D_1 when hypothesis is H_1
 - ▶ Probability is $P_d = P(D_1 \cap H_1)$

Maximum likelihood criterion

- ▶ Choose the hypothesis that **seems most likely** given the observed sample r
- ▶ The **likelihood** of an observation r = the probability density of r given a hypothesis H_0 or H_1
- ▶ Likelihood in case of hypothesis H_0 : $w(r|H_0)$
 - ▶ r is only noise, so value is taken from the noise distribution
- ▶ Likelihood in case of hypothesis H_1 : $w(r|H_1)$
 - ▶ r is $A + \text{noise}$, so value is taken from the distribution of $(A + \text{noise})$
- ▶ **Likelihood ratio** test:

$$\frac{w(r|H_1)}{w(r|H_0)} \underset{H_0}{\overset{H_1}{\gtrless}} 1$$

Graphical interpretation

- ▶ Consider noise having a normal distribution
- ▶ Plot the two density functions for H_0 , H_1

Decision via threshold

- ▶ Likelihood ratio test for ML = comparing r with a threshold T
- ▶ The threshold = the cross-over point of the two distributions

Normal noise

- ▶ Particular case: the noise has normal distribution $\mathcal{N}(0, \sigma^2)$
- ▶ Likelihood ratio is $\frac{w(r|H_1)}{r|H_0} = \frac{e^{-\frac{(r-A)^2}{2\sigma^2}}}{e^{-\frac{r^2}{2\sigma^2}}} \underset{H_0}{\overset{H_1}{\geq}} 1$
- ▶ For normal distribution, it is easier to apply *natural logarithm* to the terms
 - ▶ logarithm is a monotonic increasing function, so it won't change the comparison
 - ▶ if $A < B$, then $\log(A) < \log(B)$
- ▶ The **log-likelihood** of an observation = the logarithm of the likelihood value
 - ▶ usually the natural logarithm, but any one can be used

Log-likelihood test for ML

- ▶ For normal noise, the ML decision means the log-likelihood test

$$\frac{(r - A)^2}{r^2} \underset{H_0}{\overset{H_1}{\geq}} 1$$

- ▶ Applying square root

$$\frac{|r - A|}{|r|} \underset{H_0}{\overset{H_1}{\geq}} 1$$

- ▶ $|r - A|$ = distance from r to A , $|r|$ = distance from r to 0
- ▶ ML decision with normal noise: choose the value 0 or A which is **nearest** to r
 - ▶ very general principle, encountered in many other scenarios
 - ▶ also known as **nearest neighbor** principle / decision
 - ▶ ML receiver is also known as **minimum distance receiver**
 - ▶ equivalent with setting a threshold $T = \frac{A}{2}$

Generalizations

- ▶ What if the noise has another distribution?
 - ▶ Threshold T is still the cross-over point, whatever that is
 - ▶ There can be more cross-overs, so multiple thresholds
 - ▶ Can think that \mathbb{R} axis is split into **decision regions** R_0 and R_1
- ▶ What if the noise distributions are different for H_0 and H_1 ?
 - ▶ Threshold T is the cross-over point, whatever that is
- ▶ What if the signal $s_0(t)$ (for H_0) is not 0, but another constant value B ?
 - ▶ T is the crossover point, the distributions are centered on B and A
 - ▶ In case of normal noise, choose B or A whichever is nearest (threshold is at middle point)

Generalizations

- ▶ What if we have more than two signal levels?
 - ▶ e.g. 4 possible signals: -6, -2, 2, 6
 - ▶ Just choose the most likely hypothesis, out of 4 likelihood functions
 - ▶ Not a single threshold value, now there are more

Exercises

- ▶ A signal can have two possible values, 0 or 5. The receiver takes one sample with value $r = 2.25$
 1. Considering that the noise is white gaussian noise, what signal is decided based on the Maximum Likelihood criterion?
 2. What if the signal 0 is affected by gaussian noise $\mathcal{N}(0, 0.5)$, while the signal 5 is affected by uniform noise $\mathcal{U}[-4, 4]$?
 3. Repeat a. and b. assuming the value 0 is replaced by -1
- ▶ A signal can have four possible values: $-6, -2, 2, 6$. Each value lasts for 1 second. The signal is affected by white noise with normal distribution. The receiver takes 1 sample per second. Using ML criterion, decide what signal has been transmitted, if the received samples are:

4, 6.6, -5.2 , 1.1, 0.3, -1.5 , 7, -7 , 4.4

Computing conditional error probabilities

- ▶ We can compute the conditional probabilities of errors
- ▶ Consider the decision regions:
 - ▶ R_0 : when $r \in R_0$, decision is D_0 , i.e. $(-\infty, T)$ for gaussian noise
 - ▶ R_1 : when $r \in R_1$, decision is D_1 , i.e. $[T, \infty)$ for gaussian noise
- ▶ Probability of false alarm **if** original signal is $s_0(t)$

$$P(D_1|H_0) = \int_{R_1} w(r|H_0)dx$$

- ▶ Probability of miss **if** original signal is $s_1(t)$

$$P(D_0|H_1) = \int_{R_0} w(r|H_1)dx$$

- ▶ These probabilities do not account for the probability that the signal actually is $s_0(t)$ or $s_1(t)$
 - ▶ they are **conditional** (“if”)

Computing conditional error probabilities




Figure 2: Conditional error probabilities

[image from <http://gru.stanford.edu/doku.php/tutorials/sdt>]

Reminder: the Bayes rule

- ▶ Reminder: the Bayes rule

$$P(A \cap B) = P(B|A) \cdot P(A)$$

- ▶ Interpretation
 - ▶ The probability $P(A)$ is taken out from $P(B|A)$
 - ▶ $P(B|A)$ gives no information on $P(A)$, the chances of A actually happening
 - ▶ Example: $P(\text{score} \mid \text{shoot}) = \frac{1}{2}$. How many goals are scored?

Exercise

- ▶ A signal can have two possible values, 0 or 5. The signal 0 is affected by gaussian noise $\mathcal{N}(0, 0.5)$, while the signal 5 is affected by uniform noise $\mathcal{U}[-4, 4]$. The receiver performs ML decision based on a single sample.
 1. Compute the probability of a wrong decision when the original signal is $s_0(t)$
 2. Compute the probability of a wrong decision when the original signal is $s_1(t)$

Pitfalls of ML decision criterion

- ▶ The ML is based on comparing **conditional** probability density functions
 - ▶ conditioned by H_0 or by H_1
- ▶ Conditioning by H_0 and H_1 ignores the probability of H_0 or H_1 actually happening
 - ▶ We don't know how $p(H_0)$ or $P(H_1)$
- ▶ If $p(H_0) > p(H_1)$, we may want to move the threshold towards H_1 , and vice-versa
 - ▶ because it is more likely that the signal is $s_0(t)$
 - ▶ and thus we want to “encourage” decision D_0

The minimum error probability criterion

- ▶ Takes into account the probabilities $P(H_0)$ and $P(H_1)$
- ▶ Goal is to **minimize the total probability of error** P_e
 - ▶ errors = false alarms and misses
- ▶ We need to find the decision regions R_0 and R_1

Probability of error

- ▶ Probability of false alarm

$$\begin{aligned}P(D_1 \cap H_0) &= P(D_1|H_0) \cdot P(H_0) \\&= \int_{R_1} w(r|H_0) dx \cdot P(H_0) \\&= (1 - \int_{R_0} w(r|H_0) dx) \cdot P(H_0)\end{aligned}$$

- ▶ Probability of miss

$$\begin{aligned}P(D_0 \cap H_1) &= P(D_0|H_1) \cdot P(H_1) \\&= \int_{R_0} w(r|H_1) dx \cdot P(H_1)\end{aligned}$$

- ▶ The sum is

$$P_e = P(H_0) + \int_{R_0} [w(r|H_1) \cdot P(H_1) - w(r|H_0) \cdot P(H_0)] dx$$

Minimum probability of error

- ▶ We want to minimize P_e , i.e. to minimize the integral
- ▶ To minimize the integral, we choose R_0 such that for all $r \in R_0$, the term inside the integral is **negative**
 - ▶ because integrating over all the interval where the function is negative ensures minimum value of integral
- ▶ So, when $w(r|H_1) \cdot P(H_1) - w(r|H_0) \cdot P(H_0) < 0$ we have $r \in R_0$, i.e. decision D_0
- ▶ Conversely, when $w(r|H_1) \cdot P(H_1) - w(r|H_0) \cdot P(H_0) > 0$ we have $r \in R_1$, i.e. decision D_1
- ▶ Therefore

$$w(r|H_1) \cdot P(H_1) - w(r|H_0) \cdot P(H_0) \underset{H_0}{\overset{H_1}{\gtrless}} 0$$

$$\frac{w(r|H_1)}{w(r|H_0)} \underset{H_0}{\overset{H_1}{\gtrless}} \frac{P(H_0)}{P(H_1)}$$

Interpretation

- ▶ Similar to ML, but threshold depends on probabilities of the two hypotheses
 - ▶ When one hypotheses is more likely than the other, the threshold is pushed in its favor, towards the other
- ▶ Also based on a **likelihood ratio** test, just like ML

Minimum probability of error - gaussian noise

- Assuming the noise is gaussian (normal), $\mathcal{N}(0, \sigma^2)$

$$w(r|H_1) = e^{-\frac{(r-A)^2}{2\sigma^2}}$$

$$w(r|H_0) = e^{-\frac{r^2}{2\sigma^2}}$$

- Apply natural logarithm

$$-\frac{(r-A)^2}{2\sigma^2} + \frac{r^2}{2\sigma^2} \underset{H_0}{\overset{H_1}{\geq}} \ln \left(\frac{P(H_0)}{P(H_1)} \right)$$

- Equivalently

$$2rA - A^2 \underset{H_0}{\overset{H_1}{\geq}} 2\sigma^2 \cdot \ln \left(\frac{P(H_0)}{P(H_1)} \right)$$

$$r \underset{H_0}{\overset{H_1}{\geq}} \underbrace{\frac{A^2 + 2\sigma^2 \cdot \ln \left(\frac{P(H_0)}{P(H_1)} \right)}{2A}}_T$$

Decision regions

- ▶ We still compare with a threshold T , but its value is shifted towards the less probable hypothesis
 - ▶ T depends on the ratio $\frac{P(H_0)}{P(H_1)}$
- ▶ Decision regions
 - ▶ $R_0 = (-\infty, T]$
 - ▶ $R_1 = [T, \infty)$
 - ▶ will be different for other noise distributions (non-gaussian)

- An information source provides two messages with probabilities $p(a_0) = \frac{2}{3}$ and $p(a_1) = \frac{1}{3}$. The messages are encoded as constant signals with values -5 (a_0) and 5 (a_1). The signals are affected by gaussian noise $\mathcal{N}(0, \sigma^2 = 1)$. The receiver takes one sample r . Decision is done by comparing r with a threshold value T , as follows: if $r < T$ it is decided that the transmitted message is a_0 , otherwise it is a_1 .
1. Find the threshold value T according to the minimum probability of error criterion
 2. What if the signal 5 is affected by uniform noise $\mathcal{U}[-4, 4]$?
 3. What are the probabilities of false alarm and of miss?

Minimum risk (cost) criterion

- ▶ What if we care more about one type of errors (e.g. false alarms) than other kind (e.g. miss)?
- ▶ Minimum risk (cost) criterion: assign costs to decisions, minimize average cost
 - ▶ C_{ij} = cost of decision D_i when true hypothesis was H_j
 - ▶ C_{00} = cost for good detection D_0 in case of H_0
 - ▶ C_{10} = cost for false alarm (detection D_1 in case of H_0)
 - ▶ C_{01} = cost for miss (detection D_0 in case of H_1)
 - ▶ C_{11} = cost for good detection D_1 in case of H_1
- ▶ The risk = the average cost

$$R = C_{00}P(D_0 \cap H_0) + C_{10}P(D_1 \cap H_0) + C_{01}P(D_0 \cap H_1) + C_{11}P(D_1 \cap H_1)$$

- ▶ Minimum risk criterion: **minimize the risk R**

Computations

- ▶ Proof on table:
 - ▶ Use Bayes rule
 - ▶ Notations: $w(r|H_j)$ (*likelihood*)
 - ▶ Probabilities: $\int_{R_i} w(r|H_j) dV$
- ▶ Conclusion, **decision rule is**

$$\frac{w(r|H_1)}{w(r|H_0)} \underset{H_0}{\overset{H_1}{\gtrless}} \frac{(C_{10} - C_{00})p(H_0)}{(C_{01} - C_{11})p(H_1)}$$

Interpretation

- ▶ Similar to ML and to minimum probability of error criteria
 - ▶ also uses a **likelihood ratio** test
- ▶ Both probabilities and the assigned costs can move threshold towards one side or the other
- ▶ If $C_{10} - C_{00} = C_{01} - C_{11}$, reduces to previous criterion (minimum probability of error)
 - ▶ e.g. if $C_{00} = C_{11} = 0$, and $C_{10} = C_{01}$

In gaussian noise

- ▶ If the noise is gaussian (normal), then similar to other criteria, apply logarithm
- ▶ Equivalently

$$-(r - A)^2 + r^2 \underset{H_0}{\overset{H_1}{\geq}} \underbrace{2\sigma^2 \cdot \ln \left(\frac{(C_{10} - C_{00})p(H_0)}{(C_{01} - C_{11})p(H_1)} \right)}_C$$

$$r \underset{H_0}{\overset{H_1}{\geq}} \underbrace{\frac{A^2 + 2\sigma^2 \cdot \ln \left(\frac{(C_{10} - C_{00})p(H_0)}{(C_{01} - C_{11})p(H_1)} \right)}{2A}}_T$$

Example

- ▶ Example at blackboard: 0 / 5, random noise with $N(0, \sigma^2)$, one sample

Two non-zero levels

- ▶ What if the s_0 signal is not 0, but another constant signal $s_0 = B$?
- ▶ Noise distribution $w(r|H_0)$ is centered on B , not 0
- ▶ Otherwise, everything else stays the same
- ▶ Performance is defined by the gap between the two levels ($A - B$)
 - ▶ same performance if $s_0 = 0$, $s_1 = A$ or if $s_0 = -\frac{A}{2}$ and $s_1 = \frac{A}{2}$
- ▶ Valid for all decision criteria

Differential vs single-ended signalling

- ▶ Single-ended signaling: one signal is 0, other is non-zero
 - ▶ $s_0 = 0, s_1 = A$
- ▶ Differential signaling: use two non-zero levels with different sign, same absolute value
 - ▶ $s_0 = -\frac{A}{2}, s_1 = \frac{A}{2}$
- ▶ Which is better?

Differential vs single-ended signalling

- ▶ If gap difference between levels is the same, performance is the same
- ▶ Average power of a signal = average squared value
- ▶ For differential signal: $P = \left(\pm \frac{A}{2}\right)^2 = \frac{A^2}{4}$
- ▶ For signal ended signal: $P = P(H_0) \cdot 0 + P(H_1) (A)^2 = \frac{A^2}{2}$
 - ▶ assuming equal probabilities $P(H_0) = P(H_1) = \frac{1}{2}$
- ▶ Differential uses half the power of single-ended (i.e. better)

Summary of criteria

- ▶ We have seen decision based on 1 sample r , between 2 constant levels
- ▶ All decisions are based on a likelihood-ratio test

$$\frac{w(r|H_1)}{w(r|H_0)} \underset{H_0}{\overset{H_1}{\gtrless}} K$$

- ▶ Different criteria differ in the chosen value of K (likelihood threshold)
- ▶ Depending on the noise distributions, the real axis is partitioned into regions
 - ▶ region R_0 : if r is in here, decide D_0
 - ▶ region R_1 : if r is in here, decide D_1
 - ▶ e.g. $R_0 = (-\infty, \frac{A+B}{2}]$, $R_1 = (\frac{A+B}{2}, \infty)$ (ML)

Receiver Operating Characteristic

- ▶ The receiver performance is usually represented with “**Receiver Operating Characteristic**” (ROC) graph
- ▶ It is a graph of hit probability $P_d = P(D_1 \cap H_1)$ (correct detection) as a function of false alarm probability $P_{fa} = P(D_1 \cap H_0)$




Figure 3: Sample ROC curves

Receiver Operating Characteristic

- ▶ There is always a **tradeoff** between good P_d and bad P_{fa}
 - ▶ to increase P_d one must also increase P_{fa}
 - ▶ if we want to make sure we don't miss any real detections (increase P_d), we pay by increasing the chances of false alarms
- ▶ Different criteria = different likelihood thresholds K = different points on the graph = different tradeoffs
 - ▶ but the tradeoff cannot be avoided
- ▶ How to improve the receiver?
 - ▶ i.e. increase P_D while keeping P_{fa} the same

Performance of likelihood-ratio decoding in WGN

- ▶ WGN = “White Gaussian Noise”
- ▶ Assume equal probabilities $P(H_0) = P(H_1) = \frac{1}{2}$
- ▶ All decisions are based on a likelihood-ratio test

$$\frac{w(r|H_1)}{w(r|H_0)} \underset{H_0}{\overset{H_1}{\gtrless}} K$$

- ▶ Hit probability is

$$\begin{aligned} P_{hit} &= P(D_1|H_1)P(H_1) \\ &= P(H_1) \int_T^\infty w(r|H_1) \\ &= P(H_1)(F(\infty) - F(T)) \\ &= \frac{1}{4} \left(1 - \operatorname{erf} \left(\frac{T - A}{\sqrt{2}\sigma} \right) \right) \\ &= Q \left(\frac{T - A}{\sqrt{2}\sigma} \right) \end{aligned}$$

Performance of likelihood-ratio decoding in WGN

- False alarm probability is

$$\begin{aligned}P_{fa} &= P(D_1|H_0)P(H_0) \\&= P(H_0) \int_T^\infty w(r|H_0) \\&= P(H_0)(F(\infty) - F(T)) \\&= \frac{1}{4} \left(1 - \operatorname{erf} \left(\frac{T - 0}{\sqrt{2}\sigma} \right) \right) \\&= Q \left(\frac{T}{\sqrt{2}\sigma} \right)\end{aligned}$$

- Therefore $\frac{T}{\sqrt{2}\sigma} = Q^{-1}(P_{fa})$
- Replacing in P_{hit} yields

$$P_{hit} = Q \left(\underbrace{Q^{-1}(P_{fa})}_{\text{constant}} - \frac{A}{\sqrt{2}\sigma} \right)$$

Signal-to-noise ratio

- ▶ **Signal-to-noise ratio (SNR)** = $\frac{\text{power of original signal}}{\text{power of noise}}$
- ▶ Average power of a signal = average squared value = $\overline{X^2}$
 - ▶ Original signal power is $\frac{A^2}{2}$
 - ▶ Noise power is $\overline{X^2} = \sigma^2$ (when noise mean value $\mu = 0$)
- ▶ In our case, $\text{SNR} = \frac{A^2}{2\sigma^2}$

$$P_{hit} = Q \left(\underbrace{Q^{-1}(P_{fa})}_{\text{constant}} - \sqrt{\text{SNR}} \right)$$

- ▶ For a fixed P_{fa} , P_{hit} increases with SNR
 - ▶ Q is a monotonic decreasing function

Performance depends on SNR

- ▶ Receiver performance increases with SNR increase
 - ▶ high SNR: good performance
 - ▶ poor SNR: bad performance




Figure 4: Detection performance depends on SNR

[source: *Fundamentals of Statistical Signal Processing*, Steven Kay]