

Decision and Estimation in Data Processing

Chapter III. Elements of Estimation Theory

II.1 Introduction

What means “Estimation”?

- ▶ A sender transmits a signal $s_{\Theta}(t)$ which depends on an **unknown** parameter Θ
- ▶ The signal is affected by noise, we receive $r(t) = s_{\Theta}(t) + \text{noise}$
- ▶ We want to **find out** the correct value of the parameter
 - ▶ based on samples from the received signal, or the full continuous signal
 - ▶ available data is noisy \Rightarrow we “estimate” the parameter
- ▶ The found value is $\hat{\Theta}$, **the estimate** of Θ (“estimatul”, rom)
 - ▶ there will always be some estimation error $\epsilon = \hat{\Theta} - \Theta$
- ▶ Examples:
 - ▶ Unknown amplitude of constant signal: $r(t) = A + \text{noise}$, estimate A
 - ▶ Unknown phase of sine signal: $r(t) = \cos(2\pi ft + \phi)$, estimate ϕ
 - ▶ Record speech signal, estimate/decide what word is pronounced

Estimation vs Decision

- ▶ Consider the following estimation: $r(t) = A + \text{noise}$, estimate A
- ▶ For detection, we have to choose between **two known values** of A :
 - ▶ i.e. A can be 0 or 5 (hypotheses H_0 and H_1)
- ▶ For estimation, A can be anything \Rightarrow we choose between **infinite number of options** for A :
 - ▶ A might be any value in \mathbb{R} , in general

Estimation vs Decision

- ▶ Detection = Estimation constrained to only a few discrete options
- ▶ Estimation = Detection with an infinite number of options available
- ▶ The statistical methods used are quite similar
 - ▶ In practice, distinction between Estimation and Detections is somewhat blurred
 - ▶ (e.g. when choosing between 1000 hypotheses, do we call it “Detection” or “Estimation”?)

Available data

- ▶ The available data is the received signal $r(t)$
 - ▶ affected by noise, and depending on the unknown Θ
- ▶ We consider **N samples** from $r(t)$, taken at some sample times t_i

$$\mathbf{r} = [r_1, r_2, \dots, r_N]$$

- ▶ Each sample r_i is a random variable that depends on Θ (and the noise)
 - ▶ Each sample has a distribution that depends on Θ

$$w_i(r_i; \Theta)$$

- ▶ The whole sample vector \mathbf{r} is a N-dimensional random variable that depends on Θ (and the noise)
 - ▶ It has a N-dimensional distribution that depends on Θ

$$w(\mathbf{r}; \Theta)$$

Types of estimation

- ▶ We consider estimating a parameter Θ under two circumstances:
 1. No distribution is known about the parameter, except maybe some allowed range (e.g. $\Theta > 0$)
 - ▶ The parameter can be any value in the allowed range, equally likely
 2. We know a distribution $p(\Theta)$ for Θ , which tells us the values of Θ that are more likely than others
 - ▶ this is known as a *priori* (or *prior*) distribution (i.e. “known beforehand”)

II.2 Maximum Likelihood estimation

Maximum Likelihood definition

- ▶ When no distribution is known about the parameter, we use a method known as **Maximum Likelihood estimation (MLE)**
- ▶ The distribution of the received data, $w(\mathbf{r}; \Theta)$, is known as the **likelihood function**
 - ▶ we know the vector \mathbf{r} we received, so this is a constant
 - ▶ the unknown variable in this function is Θ

$$L(\Theta) = w(\mathbf{r}; \Theta)$$

Maximum Likelihood definition

Maximum Likelihood (ML) Estimation:

- ▶ The estimate $\hat{\Theta}$ is **the value that maximizes the likelihood of the observed data**
 - ▶ i.e. the value Θ that maximizes $L(\Theta) = w(\mathbf{r}; \Theta)$

$$\hat{\Theta} = \arg \max_{\Theta} L(\Theta) = \arg \max_{\Theta} w(\mathbf{r}; \Theta)$$

- ▶ If Θ is allowed to live only in a certain range, restrict the maximization only to that range.

How to solve

- ▶ How to solve the maximization problem?
 - ▶ i.e. how to find the estimate Θ which maximizes $L(\Theta)$
- ▶ Find maximum by setting derivative to 0

$$\frac{dL(\Theta)}{d\Theta} = 0$$

- ▶ We can also maximize **natural logarithm** of the likelihood function (“log-likelihood function”)

$$\frac{d \ln(L(\Theta))}{d\Theta} = 0$$

Solving procedure

Solving procedure:

1. Find the function

$$L(\Theta) = w(\mathbf{r}; \Theta)$$

2. Set the condition that derivative of $L(\Theta)$ or $\ln(L(\Theta))$ is 0

$$\frac{dL(\Theta)}{d\Theta} = 0, \text{ or } \frac{d \ln(L(\Theta))}{d\Theta} = 0$$

3. Solve and find the value $\hat{\Theta}$
4. Check that second derivative at point $\hat{\Theta}$ is negative, to check that point is a maximum
 - ▶ because derivative = 0 for both maximum and minimum points


Examples:

Estimating a constant signal in gaussian noise:

- ▶ Find the ML estimate of a constant value A from 5 noisy measurements $r_i = A + \text{noise}$ with values $[5, 7, 8, 6.1, 5.3]$. The noise is AWGN $\mathcal{N}(\mu = 0, \sigma^2)$.
- ▶ Solution: at whiteboard.
- ▶ The estimate \hat{A} is the average value of the samples (not surprisingly)

Numerical simulation

/home/ncleju/.local/bin/pweave:5: UserWarning: In Matplotlib individual lines on a stem plot will be added as a LineCollection instead of individual lines. This significantly improves the performance of a stem plot. To remove this warning and switch new behaviour, set the "use_line_collection" keyword argument



Curve fitting

- ▶ Estimation = curve fitting
- ▶ From the previous graphical example:
 - ▶ we have some data \mathbf{r}
 - ▶ we know the shape of the signal = a line (constant A)
 - ▶ we're fitting the best line through the data

General signal in AWGN

- ▶ Consider that the true underlying signal is $s_{\Theta}(t)$
- ▶ Consider AWGN noise $\mathcal{N}(\mu = 0, \sigma^2)$.
- ▶ The samples r_i are taken at sample moments t_i
- ▶ The samples r_i have normal distribution with average $s_{\Theta}(t_i)$ and variance σ^2
- ▶ Overall likelihood function = product of likelihoods for each sample r_i

$$\begin{aligned} L(\Theta) &= \prod_{i=1}^N \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(r_i - s_{\Theta}(t_i))^2}{2\sigma^2}} \\ &= \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{\sum (r_i - s_{\Theta}(t_i))^2}{2\sigma^2}} \end{aligned}$$

General signal in AWGN

- The log-likelihood is

$$\ln(L(\Theta)) = \underbrace{\ln\left(\frac{1}{\sigma\sqrt{2\pi}}\right)}_{\text{constant}} - \frac{\sum (r_i - s_{\Theta}(t_i))^2}{2\sigma^2}$$

General signal in AWGN

- ▶ The maximum of the function = the minimum of the exponent

$$\hat{\Theta} = \arg \max_{\Theta} w(r; \Theta) = \arg \min \sum (r_i - s_{\Theta}(t_i))^2$$

- ▶ The term $\sum (r_i - s_{\Theta}(t_i))^2$ is the **squared distance** $d(\mathbf{r}, s_{\Theta})$

$$d(\mathbf{r}, s_{\Theta}) = \sqrt{\sum (r_i - s_{\Theta}(t_i))^2}$$

$$(d(\mathbf{r}, s_{\Theta}))^2 = \sum (r_i - s_{\Theta}(t_i))^2$$

General signal in AWGN

- ▶ ML estimation can be rewritten as:

$$\hat{\Theta} = \arg \max_{\Theta} w(r; \Theta) = \arg \min d(\mathbf{r}, \mathbf{s}_{\Theta})^2$$

- ▶ ML estimate $\hat{\Theta}$ = the value that makes $s_{\Theta}(t_i)$ **closest to the received values \mathbf{r}**
 - ▶ closer = more likely
 - ▶ closest = most likely = maximum likelihood
- ▶ ML estimation = minimization of distance
- ▶ True for all kinds of vector spaces
 - ▶ vectors with N elements, continuous signals, etc
 - ▶ just change the definition of the distance function

General signal in AWGN

- Find maximum by setting derivative to 0

$$\frac{d \ln (L(\Theta))}{d \Theta} = 0$$

means

$$\sum (r_i - s_{\Theta}(t_i)) \frac{ds_{\Theta}(t_i)}{d \Theta} = 0$$


Numerical simulation

Estimating the frequency f of a cosine signal

- ▶ Find the Maximum Likelihood estimate of the frequency f of a cosine signal, from 10 noisy measurements $r_i = \cos(2\pi f t_i) + \text{noise}$ with values [...]. The noise is AWGN $\mathcal{N}(\mu = 0, \sigma^2)$. The sample times $t_i = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]$
- ▶ Solution: at whiteboard.


Numerical simulation

The likelihood function is:



Numerical simulation

```
/home/ncleju/.local/bin/pweave:15: UserWarning: In Matplotlib
individual lines on a stem plot will be added as a LineCollection
instead of individual lines. This significantly improves the
performance of a stem plot. To remove this warning and switch
new behaviour, set the "use_line_collection" keyword argument
True frequency = 0.070000, Estimate = 0.071515
```



ML Estimation and ML Detection

- ▶ In ML Estimation, the estimate $\hat{\Theta}$ is the value that maximizes the likelihood function
- ▶ In ML Detection, the decision criterion $\frac{w(r|H_1)}{w(r|H_0)} \underset{H_0}{\overset{H_1}{\gtrless}} 1$ means “choose the hypothesis that maximizes the likelihood function”.
- ▶ Therefore it is the same principle, merely in a different context:
 - ▶ in Detection we are restricted to a few predefined options
 - ▶ in Estimation we are unrestricted \Rightarrow choose the maximizing value

Loss function

- ▶ The distance $d(\mathbf{r}, \mathbf{s}_\Theta)$ is known as the “**loss function**” in machine learning terminology
 - ▶ the Euclidean distance = the “**Mean Squared Error**” (MSE) loss function
- ▶ For a given \mathbf{r} , the MSE loss = $\frac{1}{N}d(\mathbf{r}, \mathbf{s}_\Theta)$
- ▶ Other loss functions are used in different scenarios

Multiple parameters

- ▶ What if we have more than one parameter?
 - ▶ e.g. unknown parameters are the amplitude, frequency and the initial phase of a cosine:

$$s(t) = A \cos(2\pi ft + \phi)$$

- ▶ We can consider the parameter Θ to be a vector:

$$\Theta = [\Theta_1, \Theta_2, \dots, \Theta_M]$$

- ▶ e.g. $\Theta = [\Theta_1, \Theta_2, \Theta_3] = [A, f, \phi]$

Gradient Descent

- ▶ How to estimate the parameters Θ in complicated cases?
 - ▶ e.g. in real life applications
 - ▶ usually there are many parameters (Θ is a vector)
- ▶ Typically it is impossible to get the optimal values directly
- ▶ Improve them iteratively with **Gradient Descent** algorithm or its variations

Gradient Descent procedure

1. Start with some random parameter values $\Theta^{(0)}$
2. Repeat for each iteration k :
 - 2.1 Compute loss value $L(\Theta^{(k)})$
 - 2.2 Compute derivative $\frac{\partial L}{\partial \Theta_i^{(k)}}$ for each Θ_i
 - 2.3 Update all values Θ_i by subtracting the derivative

$$\Theta_i^{(k+1)} = \Theta_i^{(k)} - \mu \frac{\partial L}{\partial \Theta_i^{(k)}}$$

► or, in vector form:

$$\Theta^{(k+1)} = \Theta^k - \mu \frac{\partial L}{\partial \Theta^{(k)}}$$

3. Until termination criterion (e.g. parameters don't change much)

Gradient Descent explained

- ▶ Explanations at blackboard
- ▶ Simple example: logistic regression on 2D-data
 - ▶ maybe do example at blackboard

Neural Networks

- ▶ The most prominent example is **Artificial Neural Networks** (a.k.a. Neural Networks, Deep Learning, etc.)
 - ▶ Can be regarded as ML estimation
 - ▶ Use loss function (typically not MSE, but others)
 - ▶ Use Gradient Descent to update parameters
 - ▶ State-of-the-art applications: image classification/recognition, automated driving etc.
- ▶ More info on neural networks / machine learning:
 - ▶ look up online courses, books (e.g. prof. Iulian Ciocoiu's book)
 - ▶ join the IASI AI Meetup

II.3 Bayesian estimation

Prior distribution

- ▶ Suppose we know beforehand a distribution of Θ , $w(\Theta)$
 - ▶ we know beforehand how likely it is to have a certain value
 - ▶ known as *a priori* distribution or *prior* distribution
- ▶ The estimation must take it into account
 - ▶ the estimate will be slightly “moved” towards more likely values
- ▶ Known as “Bayesian estimation”
 - ▶ Thomas Bayes = discovered the Bayes rule
 - ▶ Stuff related to Bayes rule are often named “Bayesian”

Cost function

- ▶ The **estimation error** is the difference between the estimate $\hat{\Theta}$ and the true value Θ

$$\epsilon = \hat{\Theta} - \Theta$$

- ▶ The **cost function** $C(\epsilon)$ assigns a cost to each possible estimation error
 - ▶ when $\epsilon = 0$, the cost $C(0) = 0$
 - ▶ small errors ϵ have small costs
 - ▶ large errors ϵ have large costs
- ▶ Usual types of cost functions:
 - ▶ Quadratic: $C(\epsilon) = \epsilon^2 = (\hat{\Theta} - \Theta)^2$
 - ▶ Uniform ("hit or miss"): $C(\epsilon) = \begin{cases} 0, & \text{if } |\epsilon| = |\hat{\Theta} - \Theta| \leq E \\ 1, & \text{if } |\epsilon| = |\hat{\Theta} - \Theta| > E \end{cases}$
 - ▶ Linear: $C(\epsilon) = |\epsilon| = |\hat{\Theta} - \Theta|$
 - ▶ draw them at whiteboard

The Bayesian risk

- ▶ For each pair of values \mathbf{r} and Θ , $w(\mathbf{r}; \Theta)$ tells us how likely it is to have them
- ▶ Multiplying with $C(\epsilon)$ gives us the cost, for each \mathbf{r} and Θ

$$C(\epsilon)w(\mathbf{r}; \Theta)$$

- ▶ Integrating over Θ gives the cost for a certain \mathbf{r}

$$\int_{-\infty}^{\infty} C(\epsilon)w(\mathbf{r}; \Theta)d\Theta$$

- ▶ Further integrating also over \mathbf{r} gives the global cost for all \mathbf{r} and all Θ

$$R = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} C(\epsilon)w(\mathbf{r}; \Theta)d\Theta d\mathbf{r}$$

Minimizing the risk

- ▶ We want to minimize the risk R
- ▶ Bayes rule: $w(\mathbf{r}; \Theta) = w(\Theta|\mathbf{r})w(\mathbf{r})$
- ▶ Replacing in R , we obtain

$$\begin{aligned} R &= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} C(\epsilon) w(\Theta|\mathbf{r}) w(\mathbf{r}) d\Theta d\mathbf{r} \\ &= \int_{-\infty}^{\infty} w(\mathbf{r}) \left[\int_{-\infty}^{\infty} C(\epsilon) w(\Theta|\mathbf{r}) d\Theta \right] d\mathbf{r} \end{aligned}$$

- ▶ Since $w(\mathbf{r}) \geq 0$, minimizing the inner integral will minimize R

$$I = \int_{-\infty}^{\infty} C(\epsilon) w(\Theta|\mathbf{r}) d\Theta$$

- ▶ Next, we'll replace $C(\epsilon)$ with its definition and derivate over $\hat{\Theta}$
 - ▶ Attention: $\hat{\Theta}$, not Θ !

MMSE estimator

- ▶ When the cost function is quadratic $C(\epsilon) = \epsilon^2 = (\hat{\Theta} - \Theta)^2$

$$I = \int_{-\infty}^{\infty} (\hat{\Theta} - \Theta)^2 w(\Theta|\mathbf{r}) d\Theta$$

- ▶ We want the $\hat{\Theta}$ that minimizes I , so we derivate

$$\frac{dI}{d\hat{\Theta}} = 2 \int_{-\infty}^{\infty} (\hat{\Theta} - \Theta) w(\Theta|\mathbf{r}) d\Theta = 0$$

- ▶ Equivalent to

$$\hat{\Theta} \underbrace{\int_{-\infty}^{\infty} w(\Theta|\mathbf{r}) d\Theta}_1 = \int_{-\infty}^{\infty} \Theta w(\Theta|\mathbf{r}) d\Theta$$

- ▶ The **Minimum Mean Squared Error (MMSE)** estimator is

$$\hat{\Theta} = \int_{-\infty}^{\infty} \Theta \cdot w(\Theta|\mathbf{r}) d\Theta$$

Interpretation

- ▶ $w(\Theta|\mathbf{r})$ is the **posterior** (or **a posteriori**) distribution
 - ▶ it is the distribution of Θ after we know the data we received
 - ▶ the prior distribution $w(\Theta)$ is the one before knowing any data
- ▶ The MMSE estimation is the **average value** of the posterior distribution

The MAP estimator

- ▶ When the cost function is uniform

$$C(\epsilon) = \begin{cases} 0, & \text{if } |\epsilon| = |\hat{\Theta} - \Theta| \leq E \\ 1, & \text{if } |\epsilon| = |\hat{\Theta} - \Theta| > E \end{cases} \quad \begin{matrix} \\ \end{matrix}$$

- Keep in mind that $\Theta = \hat{\Theta} - \epsilon$

- We obtain

$$I = \int_{-\infty}^{\hat{\Theta}-E} w(\Theta|\mathbf{r})d\Theta + \int_{\hat{\Theta}+E}^{\infty} w(\Theta|\mathbf{r})d\Theta$$

$$I = 1 - \int_{\hat{\Theta}-E}^{\hat{\Theta}+E} w(\Theta|\mathbf{r})d\Theta$$

The MAP estimator

- ▶ To minimize I , we must maximize $\int_{\hat{\Theta}-E}^{\hat{\Theta}+E} w(\Theta|\mathbf{r})d\Theta$, the integral around point $\hat{\Theta}$
- ▶ For E a very small, the function $w(\Theta|\mathbf{r})$ is approximately constant, so we pick the point where the function is maximum
- ▶ The **Maximum A Posteriori (MAP)** estimator is

$$\hat{\Theta} = \arg \max w(\Theta|\mathbf{r})$$

- ▶ $\arg \max$ = “the value which maximizes the function”
 - ▶ $\max f(x)$ = the maximum value of a function
 - ▶ $\arg \max f(x)$ = the x for which the function reaches its maximum

Interpretation

- ▶ The MAP estimator chooses Θ as the value where the posterior distribution is maximum
- ▶ The MMSE estimator chooses Θ as average value of the posterior distribution




Figure 1: MAP vs MMSE estimators

Finding the posterior distribution

- ▶ That's cool, but how do we find this posterior distribution $w(\Theta|\mathbf{r})$?
- ▶ Use the Bayes rule

$$w(\Theta|\mathbf{r}) = \frac{w(\mathbf{r}; \Theta)}{w(\mathbf{r})} = \frac{w(\mathbf{r}|\Theta) \cdot w(\Theta)}{w(\mathbf{r})}$$

- ▶ Since $w(\mathbf{r})$ is constant for a given \mathbf{r} the MAP estimator is

$$\hat{\Theta} = \arg \max w(\Theta|\mathbf{r}) = \arg \max w(\mathbf{r}|\Theta)w(\Theta)$$

- ▶ The MAP estimator is the one which **maximizes** the likelihood of the observed data, **but multiplying with the prior distribution** $w(\Theta)$
- ▶ The MMSE estimator is the **average** of the same thing

Relation with Maximum Likelihood Estimator

- ▶ The ML estimator was just $\arg \max w(\mathbf{r}|\Theta)$
- ▶ The MAP estimator = like the ML estimator but multiplied with the prior distribution $w(\Theta)$
- ▶ If $w(\Theta)$ is a constant, the MAP estimator reduces to ML
 - ▶ $w(\Theta) = \text{constant}$ means all values Θ are equally likely
 - ▶ i.e. we don't have a clue where the real Θ might be
- ▶ The MMSE estimator = like MAP, but don't take the *argmax* of the function, but its average value

Relation with Detection

- ▶ The minimum probability of error criterion $\frac{w(r|H_1)}{w(r|H_0)} \underset{H_0}{\overset{H_1}{\gtrless}} \frac{P(H_0)}{P(H_1)}$
- ▶ It can be rewritten as $w(r|H_1) \cdot P(H_1) \underset{H_0}{\overset{H_1}{\gtrless}} w(r|H_0)P(H_0)$
 - ▶ i.e. choose the hypothesis where $w(r|H) \cdot P(H)$ is maximum
 - ▶ $w(r|H_1)$, $w(r|H_0)$ are the likelihood of observed data
 - ▶ $P(H_1)$, $P(H_0)$ are the prior probabilities (known beforehand)
- ▶ The MAP estimator is where $w(\mathbf{r}|\Theta)w(\Theta)$ is maximum
 - ▶ $w(\mathbf{r}|\Theta)$ is the likelihood of observed data
 - ▶ $w(\Theta)$ is the prior distribution (known beforehand)
- ▶ Therefore it is the same principle, merely in a different context:
 - ▶ in Detection we are restricted to a few predefined options
 - ▶ in Estimation we are unrestricted \Rightarrow choose the maximizing value of the whole function

- ▶ Chapter ends here for 2018-2019 exam. Following slides not needed.

Exercise

Exercise: constant value, 3 measurement, Gaussian same σ

- ▶ We want to estimate today's temperature in Sahara
- ▶ Our thermometer reads 40 degrees, but the value was affected by Gaussian noise $\mathcal{N}(0, \sigma^2 = 2)$ (crappy thermometer)
- ▶ We know that this time of the year, the temperature is around 35 degrees, with a Gaussian distribution $\mathcal{N}(35, \sigma^2 = 2)$.
- ▶ Estimate the true temperature using ML, MAP and MMSE estimators

Exercise

Exercise: constant value, 3 measurements, Gaussian same σ

- ▶ What if he have three thermometers, showing 40, 38, 41 degrees

Exercise: constant value, 3 measurements, Gaussian different σ

- ▶ What if the temperature this time of the year has Gaussian distribution $\mathcal{N}(35, \sigma_2^2 = 3)$
 - ▶ different variance, $\sigma_2 \neq \sigma$

General signal in AWGN

- ▶ Consider that the true underlying signal is $s_{\Theta}(t)$
- ▶ Consider AWGN noise $\mathcal{N}(\mu = 0, \sigma^2)$.
- ▶ As in Maximum Likelihood function, overall likelihood function

$$w(\mathbf{r}|\Theta) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{\sum (r_i - s_{\Theta}(t_i))^2}{2\sigma^2}}$$

- ▶ But now this function is also **multiplied with** $w(\Theta)$

$$w(\mathbf{r}|\Theta) \cdot w(\Theta)$$

General signal in AWGN

- ▶ MAP estimator is the argument that maximizes this product

$$\hat{\Theta}_{MAP} = \arg \max w(\mathbf{r}|\Theta)w(\Theta)$$

- ▶ Taking logarithm

$$\begin{aligned}\hat{\Theta}_{MAP} &= \arg \max \ln (w(\mathbf{r}|\Theta)) + \ln (w(\Theta)) \\ &= \arg \max -\frac{\sum (r_i - s_{\Theta}(t_i))^2}{2\sigma^2} + \ln (w(\Theta))\end{aligned}$$

Gaussian prior

- ▶ If the prior distribution is also Gaussian $\mathcal{N}(\mu_{\Theta}, \sigma_{\Theta}^2)$

$$\ln(w(\Theta)) = -\frac{\sum(\Theta - \mu_{\Theta})^2}{2\sigma_{\Theta}^2}$$

- ▶ MAP estimation becomes

$$\hat{\Theta}_{MAP} = \arg \min \frac{\sum(r_i - s_{\Theta}(t_i))^2}{2\sigma^2} + \frac{\sum(\Theta - \mu_{\Theta})^2}{2\sigma_{\Theta}^2}$$

- ▶ Can be rewritten as

$$\hat{\Theta}_{MAP} = \arg \min d(\mathbf{r}, s_{\Theta})^2 + \underbrace{\frac{\sigma^2}{\sigma_{\Theta}^2}}_{\lambda} \cdot d(\Theta, \mu_{\Theta})^2$$

Interpretation

- ▶ MAP estimator with Gaussian noise and Gaussian prior

$$\hat{\Theta}_{MAP} = \arg \min d(\mathbf{r}, s_{\Theta})^2 + \underbrace{\frac{\sigma^2}{\sigma_{\Theta}^2}}_{\lambda} \cdot d(\Theta, \mu_{\Theta})^2$$

- ▶ $\hat{\Theta}_{MAP}$ is close to its expected value μ_{Θ} and it makes the true signal close to received data \mathbf{r}
 - ▶ Example: “search for a house that is close to job and close to the Mall”
 - ▶ λ controls the relative importance of the two terms
- ▶ Particular cases
 - ▶ σ_{Θ} very small = the prior is very specific (narrow) = λ large = second term very important = $\hat{\Theta}_{MAP}$ close to μ_{Θ}
 - ▶ σ_{Θ} very large = the prior is very unspecific = λ small = first term very important = $\hat{\Theta}_{MAP}$ close to ML estimation

Applications

- ▶ In general, practical applications:
 - ▶ can use various prior distributions
 - ▶ estimate **multiple parameters** (a vector of parameters)
- ▶ Applications
 - ▶ denoising of signals
 - ▶ signal restoration
 - ▶ signal compression

Estimator bias

- ▶ How good is an estimator?
 - ▶ Many ways to characterize
- ▶ An estimator $\hat{\Theta}$ is a **random variable**
 - ▶ can have different values, because it is computed based on the received samples, which depend on noise
 - ▶ example: in lab, try on multiple computers => slightly different results
- ▶ As a random variable, it has:
 - ▶ an average value (expected value): $E \{ \hat{\Theta} \}$
 - ▶ a variance: $E \{ (\hat{\Theta} - \Theta)^2 \}$

Estimator bias

- ▶ **Unbiased** estimator = if the average value of the estimator is the true value of Θ

$$E \{ \hat{\Theta} \} = \Theta$$

- ▶ **Biased** estimator = if the average value of the estimator is different from the true value Θ
 - ▶ the difference $E \{ \hat{\Theta} \} - \Theta$ is called **the bias** of the estimator

Estimator bias

- ▶ Example: for constant signal A with AWGN noise (zero-mean), ML estimator is $\hat{A}_{ML} = \frac{1}{N} \sum_i r_i$
- ▶ Then:

$$\begin{aligned} E \{ \hat{A}_{ML} \} &= \frac{1}{N} E \left\{ \sum_i r_i \right\} \\ &= \frac{1}{N} \sum_{i=1}^N E \{ r_i \} \\ &= \frac{1}{N} \sum_{i=1}^N E \{ A + \text{noise} \} \\ &= \frac{1}{N} \sum_{i=1}^N A \\ &= A \end{aligned}$$

- ▶ This estimator is unbiased

Estimator variance

- ▶ Unbiased estimators are good, but if the **variance** of the estimator is large, then estimated values can be far from the true value
- ▶ We prefer estimators with **small variance**, even if maybe slightly biased