

## Chapter III. Elements of Estimation Theory

## II.1 Introduction

# What means “Estimation”?

- ▶ A sender transmits a signal  $s_{\Theta}(t)$  which depends on an **unknown** parameter  $\Theta$
- ▶ The signal is affected by noise, we receive  $r(t) = s_{\Theta}(t) + \text{noise}$
- ▶ We want to **find out** the correct value of the parameter
  - ▶ based on samples from the received signal, or the full continuous signal
  - ▶ available data is noisy  $\Rightarrow$  we “estimate” the parameter
- ▶ The found value is  $\hat{\Theta}$ , **the estimate** of  $\Theta$  (“estimatul”, rom)
  - ▶ there will always be some estimation error  $\epsilon = \hat{\Theta} - \Theta$
- ▶ Examples:
  - ▶ Unknown amplitude of constant signal:  $r(t) = A + \text{noise}$ , estimate  $A$
  - ▶ Unknown phase of sine signal:  $r(t) = \cos(2\pi ft + \phi)$ , estimate  $\phi$
  - ▶ Record speech signal, estimate/decide what word is pronounced

# Estimation vs Decision

- ▶ Consider the following estimation:  $r(t) = A + \text{noise}$ , estimate  $A$
- ▶ For detection, we have to choose between **two known values** of  $A$ :
  - ▶ i.e.  $A$  can be 0 or 5 (hypotheses  $H_0$  and  $H_1$ )
- ▶ For estimation,  $A$  can be anything  $\Rightarrow$  we choose between **infinite number of options** for  $A$ :
  - ▶  $A$  might be any value in  $\mathbb{R}$ , in general

# Estimation vs Decision

- ▶ Detection = Estimation constrained to only a few discrete options
- ▶ Estimation = Detection with an infinite number of options available
- ▶ The statistical methods used are quite similar
  - ▶ In practice, distinction between Estimation and Detections is somewhat blurred
  - ▶ (e.g. when choosing between 1000 hypotheses, do we call it “Detection” or “Estimation”?)

# Available data

- ▶ The available data is the received signal  $r(t)$ 
  - ▶ affected by noise, and depending on the unknown  $\Theta$
- ▶ We consider **N samples** from  $r(t)$ , taken at some sample times  $t_i$

$$\mathbf{r} = [r_1, r_2, \dots, r_N]$$

- ▶ Each sample  $r_i$  is a random variable that depends on  $\Theta$  (and the noise)
  - ▶ Each sample has a distribution that depends on  $\Theta$

$$w_i(r_i; \Theta)$$

- ▶ The whole sample vector  $\mathbf{r}$  is a N-dimensional random variable that depends on  $\Theta$  (and the noise)
  - ▶ It has a N-dimensional distribution that depends on  $\Theta$

$$w(\mathbf{r}; \Theta)$$

# Types of estimation

- ▶ We consider estimating a parameter  $\Theta$  under two circumstances:
  1. No distribution is known about the parameter, except maybe some allowed range (e.g.  $\Theta > 0$ )
    - ▶ The parameter can be any value in the allowed range, equally likely
  2. We know a distribution  $p(\Theta)$  for  $\Theta$ , which tells us the values of  $\Theta$  that are more likely than others
    - ▶ this is known as *a priori* (or *prior*) distribution (i.e. “known beforehand”)

## II.2 Maximum Likelihood estimation



# Maximum Likelihood definition

- ▶ When no distribution is known about the parameter, we use a method known as Maximum Likelihood Estimation (MLE)
- ▶ The distribution of the received data,  $w(\mathbf{r}; \Theta)$ , is known as the **likelihood function**
  - ▶ we know the vector  $\mathbf{r}$  we received, so this is a constant
  - ▶ the unknown variable in this function is  $\Theta$

$$L(\Theta) = w(\mathbf{r}; \Theta)$$

- ▶ Maximum Likelihood Estimation: The estimate  $\hat{\Theta}$  is **the value that maximizes the likelihood of the observed data**
  - ▶ i.e. the value  $\Theta$  that maximizes  $w(r; \Theta)$

$$\hat{\Theta} = \arg \max_{\Theta} L(\Theta) = \arg \max_{\Theta} w(r; \Theta)$$

- ▶ If  $\Theta$  is allowed to live only in a certain range, restrict the maximization only to that range.

# Computations

- ▶ Find maximum by setting derivative to 0

$$\frac{dL(\Theta)}{d\Theta} = 0$$

- ▶ We can also maximize **natural logarithm** of the likelihood function (“log-likelihood function”)

$$\frac{d \ln (L(\Theta))}{d\Theta} = 0$$

# Computations

Method:

1. Find the function

$$L(\Theta) = w(\mathbf{r}; \Theta)$$

2. Set the condition that derivative of  $L(\Theta)$  or  $\ln(L(\Theta))$  is 0

$$\frac{dL(\Theta)}{d\Theta} = 0, \text{ or } \frac{d \ln(L(\Theta))}{d\Theta} = 0$$


3. Solve and find the value  $\hat{\Theta}$
4. Check that second derivative at point  $\hat{\Theta}$  is negative, to check that point is a maximum
  - ▶ because derivative = 0 for both maximum and minimum points

## Examples:

Estimating a constant signal in gaussian noise:

- ▶ Find the Maximum Likelihood estimate of a constant value  $A$  from 5 noisy measurements  $r_i = A + \text{noise}$  with values  $[5, 7, 8, 6.1, 5.3]$ . The noise is AWGN  $\mathcal{N}(\mu = 0, \sigma^2)$ .
- ▶ Solution: at whiteboard.
- ▶ The estimate  $\hat{A}$  is the average value of the samples (not surprisingly)

# Numerical simulation



# General signal in AWGN

- ▶ Consider that the true underlying signal is  $s_{\Theta}(t)$
- ▶ Consider AWGN noise  $\mathcal{N}(\mu = 0, \sigma^2)$ .
- ▶ The samples  $r_i$  are taken at sample moments  $t_i$
- ▶ The samples  $r_i$  have normal distribution with average  $s_{\Theta}(t_i)$  and variance  $\sigma^2$
- ▶ Overall likelihood function = product of likelihoods for each sample  $r_i$

$$\begin{aligned} L(\Theta) &= \prod_{i=1}^N \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(r_i - s_{\Theta}(t_i))^2}{2\sigma^2}} \\ &= \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{\sum (r_i - s_{\Theta}(t_i))^2}{2\sigma^2}} \end{aligned}$$

- ▶ The log-likelihood is

$$\ln(L(\Theta)) = \underbrace{\ln\left(\frac{1}{\sigma\sqrt{2\pi}}\right)}_{\text{constant}} - \frac{\sum (r_i - s_{\Theta}(t_i))^2}{2\sigma^2}$$

# General signal in AWGN

- ▶ The maximum of the function = the minimum of the exponent

$$\hat{\Theta} = \arg \max_{\Theta} w(r; \Theta) = \arg \min \sum (r_i - s_{\Theta}(t_i))^2$$

- ▶ The term  $\sum (r_i - s_{\Theta}(t_i))^2$  is the **squared distance**  $d(\mathbf{r}, s_{\Theta})$

$$d(\mathbf{r}, s_{\Theta}) = \sqrt{\sum (r_i - s_{\Theta}(t_i))^2}$$

$$(d(\mathbf{r}, s_{\Theta}))^2 = \sum (r_i - s_{\Theta}(t_i))^2$$

- ▶ Maximum Likelihood estimate  $\hat{\Theta}$  = the value that makes  $s_{\Theta}(t_i)$  **closest to the received values  $\mathbf{r}$** 
  - ▶ closer = more likely
  - ▶ closest = most likely = maximum likelihood
- ▶ For continuous signals? Same, but use distance function for continuous signals

# General signal in AWGN

- Find maximum by setting derivative to 0

$$\frac{d \ln(L(\Theta))}{d\Theta} = 0$$

means

$$\sum (r_i - s_{\Theta}(t_i)) \frac{ds_{\Theta}(t_i)}{d\Theta} = 0$$




## Estimating the frequency $f$ of a cosine signal

- ▶ Find the Maximum Likelihood estimate of the frequency  $f$  of a cosine signal, from 10 noisy measurements  $r_i = \cos(2\pi f t_i) + \text{noise}$  with values [...]. The noise is AWGN  $\mathcal{N}(\mu = 0, \sigma^2)$ . The sample times  $t_i = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]$
- ▶ Solution: at whiteboard.


# Numerical simulation

The likelihood function is:



# Numerical simulation

True frequency = 0.070000, Estimate = 0.071515



# ML Estimation and ML Detection

- ▶ In ML Estimation, the estimate  $\hat{\Theta}$  is the value that maximizes the likelihood function
- ▶ In ML Detection, the decision criterion  $\frac{w(r|H_1)}{w(r|H_0)} \underset{H_0}{\overset{H_1}{\geq}} 1$  means “choose the hypothesis that maximizes the likelihood function”.
- ▶ Therefore it is the same principle, merely in a different context:
  - ▶ in Detection we are restricted to a few predefined options
  - ▶ in Estimation we are unrestricted  $\Rightarrow$  choose the maximizing value

## II.3 Bayesian estimation

# Prior distribution

- ▶ Suppose we know beforehand a distribution of  $\Theta$ ,  $w(\Theta)$ 
  - ▶ we know beforehand how likely it is to have a certain value
  - ▶ known as *a priori* distribution or *prior* distribution
- ▶ The estimation must take it into account
  - ▶ the estimate will be slightly “moved” towards more likely values
- ▶ Known as “Bayesian estimation”
  - ▶ Thomas Bayes = discovered the Bayes rule
  - ▶ Stuff related to Bayes rule are often named “Bayesian”

# Cost function

- ▶ The **estimation error** is the difference between the estimate  $\hat{\Theta}$  and the true value  $\Theta$

$$\epsilon = \hat{\Theta} - \Theta$$

- ▶ The **cost function**  $C(\epsilon)$  assigns a cost to each possible estimation error
  - ▶ when  $\epsilon = 0$ , the cost  $C(0) = 0$
  - ▶ small errors  $\epsilon$  have small costs
  - ▶ large errors  $\epsilon$  have large costs
- ▶ Usual types of cost functions:
  - ▶ Quadratic:  $C(\epsilon) = \epsilon^2 = (\hat{\Theta} - \Theta)^2$
  - ▶ Uniform (“hit or miss”):  $C(\epsilon) = \begin{cases} 0, & \text{if } |\epsilon| = |\hat{\Theta} - \Theta| \leq E \\ 1, & \text{if } |\epsilon| = |\hat{\Theta} - \Theta| > E \end{cases}$
  - ▶ Linear:  $C(\epsilon) = |\epsilon| = |\hat{\Theta} - \Theta|$
  - ▶ draw them at whiteboard

# The Bayesian risk

- ▶ For each pair of values  $\mathbf{r}$  and  $\Theta$ ,  $w(\mathbf{r}; \Theta)$  tells us how likely it is to have them
- ▶ Multiplying with  $C(\epsilon)$  gives us the cost, for each  $\mathbf{r}$  and  $\Theta$

$$C(\epsilon)w(\mathbf{r}; \Theta)$$

- ▶ Integrating over  $\Theta$  gives the cost for a certain  $\mathbf{r}$

$$\int_{-\infty}^{\infty} C(\epsilon)w(\mathbf{r}; \Theta)d\Theta$$

- ▶ Further integrating also over  $\mathbf{r}$  gives the global cost for all  $\mathbf{r}$  and all  $\Theta$

$$R = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} C(\epsilon)w(\mathbf{r}; \Theta)d\Theta d\mathbf{r}$$



# Minimizing the risk

- ▶ We want to minimize the risk  $R$
- ▶ Bayes rule:  $w(\mathbf{r}; \Theta) = w(\Theta|\mathbf{r})w(\mathbf{r})$
- ▶ Replacing in  $R$ , we obtain

$$\begin{aligned} R &= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} C(\epsilon) w(\Theta|\mathbf{r}) w(\mathbf{r}) d\Theta d\mathbf{r} \\ &= \int_{-\infty}^{\infty} w(\mathbf{r}) \left[ \int_{-\infty}^{\infty} C(\epsilon) w(\Theta|\mathbf{r}) d\Theta \right] d\mathbf{r} \end{aligned}$$

- ▶ Since  $w(\mathbf{r}) \geq 0$ , minimizing the inner integral will minimize  $R$

$$I = \int_{-\infty}^{\infty} C(\epsilon) w(\Theta|\mathbf{r}) d\Theta$$

- ▶ Next, we'll replace  $C(\epsilon)$  with its definition and derivate over  $\hat{\Theta}$ 
  - ▶ Attention:  $\hat{\Theta}$ , not  $\Theta$ !

# MMSE estimator

- ▶ When the cost function is quadratic  $C(\epsilon) = \epsilon^2 = (\hat{\Theta} - \Theta)^2$

$$I = \int_{-\infty}^{\infty} (\hat{\Theta} - \Theta)^2 w(\Theta|\mathbf{r}) d\Theta$$

- ▶ We want the  $\hat{\Theta}$  that minimizes  $I$ , so we derivate

$$\frac{dI}{d\hat{\Theta}} = 2 \int_{-\infty}^{\infty} (\hat{\Theta} - \Theta) w(\Theta|\mathbf{r}) d\Theta = 0$$

- ▶ Equivalent to

$$\hat{\Theta} \underbrace{\int_{-\infty}^{\infty} w(\Theta|\mathbf{r}) d\Theta}_1 = \int_{-\infty}^{\infty} \Theta w(\Theta|\mathbf{r}) d\Theta$$

- ▶ The **Minimum Mean Squared Error (MMSE)** estimator is

$$\hat{\Theta} = \int_{-\infty}^{\infty} \Theta \cdot w(\Theta|\mathbf{r}) d\Theta$$

# Interpretation

- ▶  $w(\Theta|\mathbf{r})$  is the **posterior** ( or **a posteriori**) distribution
  - ▶ it is the distribution of  $\Theta$  after we know the data we received
  - ▶ the prior distribution  $w(\Theta)$  is the one before knowing any data
- ▶ The MMSE estimation is the **average value** of the posterior distribution

# The MAP estimator

- ▶ When the cost function is uniform

$$C(\epsilon) = \begin{cases} 0, & \text{if } |\epsilon| = |\hat{\Theta} - \Theta| \leq E \\ 1, & \text{if } |\epsilon| = |\hat{\Theta} - \Theta| > E \end{cases} \quad \begin{matrix} \\ \end{matrix}$$

- ▶ Keep in mind that  $\Theta = \hat{\Theta} - \epsilon$
- ▶ We obtain

$$I = \int_{-\infty}^{\hat{\Theta}-E} w(\Theta|\mathbf{r})d\Theta + \int_{\hat{\Theta}+E}^{\infty} w(\Theta|\mathbf{r})d\Theta$$

$$I = 1 - \int_{\hat{\Theta}-E}^{\hat{\Theta}+E} w(\Theta|\mathbf{r})d\Theta$$

# The MAP estimator

- ▶ To minimize  $I$ , we must maximize  $\int_{\hat{\Theta}-E}^{\hat{\Theta}+E} w(\Theta|\mathbf{r})d\Theta$ , the integral around point  $\hat{\Theta}$
- ▶ For  $E$  a very small, the function  $w(\Theta|\mathbf{r})$  is approximately constant, so we pick the point where the function is maximum
- ▶ The **Maximum A Posteriori (MAP)** estimator is

$$\hat{\Theta} = \arg \max w(\Theta|\mathbf{r})$$

- ▶  $\arg \max$  = “the value which maximizes the function”
  - ▶  $\max f(x)$  = the maximum value of a function
  - ▶  $\arg \max f(x)$  = the  $x$  for which the function reaches its maximum

# Interpretation

- ▶ The MAP estimator chooses  $\Theta$  as the value where the posterior distribution is maximum
- ▶ The MMSE estimator chooses  $\Theta$  as average value of the posterior distribution

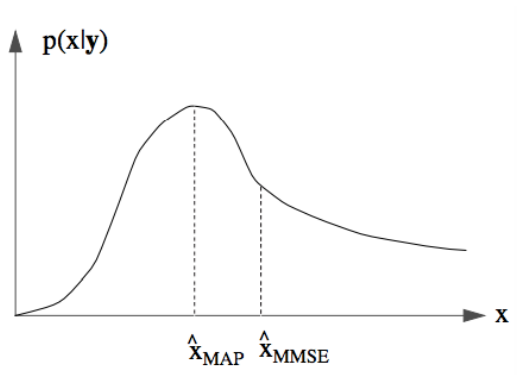


Figure 1: MAP vs MMSE estimators

# Finding the posterior distribution

- ▶ That's cool, but how do we find this posterior distribution  $w(\Theta|\mathbf{r})$ ?
- ▶ Use the Bayes rule

$$w(\Theta|\mathbf{r}) = \frac{w(\mathbf{r}; \Theta)}{w(\mathbf{r})} = \frac{w(\mathbf{r}|\Theta) \cdot w(\Theta)}{w(\mathbf{r})}$$

- ▶ Since  $w(\mathbf{r})$  is constant for a given  $\mathbf{r}$  the MAP estimator is

$$\hat{\Theta} = \arg \max w(\Theta|\mathbf{r}) = \arg \max w(\mathbf{r}|\Theta)w(\Theta)$$

- ▶ The MAP estimator is the one which **maximizes** the likelihood of the observed data, **but multiplying with the prior distribution**  $w(\Theta)$
- ▶ The MMSE estimator is the **average** of the same thing

# Relation with Maximum Likelihood Estimator

- ▶ The MLE estimator was just  $\arg \max w(\mathbf{r}|\Theta)$
- ▶ The MAP estimator = like the MLE estimator but with the prior distribution  $w(\Theta)$
- ▶ If  $w(\Theta)$  is a constant, the MAP estimator reduces to MLE
  - ▶  $w(\Theta) = \text{constant}$  means all values  $\Theta$  are equally likely
  - ▶ i.e. we don't have a clue where the real  $\Theta$  might be



# Relation with Detection

- ▶ The minimum probability of error criterion  $\frac{w(r|H_1)}{w(r|H_0)} \underset{H_0}{\overset{H_1}{\gtrless}} \frac{P(H_0)}{P(H_1)}$
- ▶ It can be rewritten as  $w(r|H_1) \cdot P(H_1) \underset{H_0}{\overset{H_1}{\gtrless}} w(r|H_0)P(H_0)$ 
  - ▶ i.e. choose the hypothesis where  $w(r|H) \cdot P(H)$  is maximum
  - ▶  $w(r|H_1)$ ,  $w(r|H_0)$  are the likelihood of observed data
  - ▶  $P(H_1)$ ,  $P(H_0)$  are the prior probabilities (known beforehand)
- ▶ The MAP estimator is where  $w(\mathbf{r}|\Theta)w(\Theta)$  is maximum
  - ▶  $w(\mathbf{r}|\Theta)$  is the likelihood of observed data
  - ▶  $w(\Theta)$  is the prior distribution (known beforehand)
- ▶ Therefore it is the same principle, merely in a different context:
  - ▶ in Detection we are restricted to a few predefined options
  - ▶ in Estimation we are unrestricted  $\Rightarrow$  choose the maximizing value of the whole function

# Exercise

Exercise: constant value, 3 measurement, Gaussian same  $\sigma$

- ▶ We want to estimate today's temperature in Sahara
- ▶ Our thermometer reads 40 degrees, but the value was affected by Gaussian noise  $\mathcal{N}(0, \sigma^2 = 2)$  (crappy thermometer)
- ▶ We know that this time of the year, the temperature is around 35 degrees, with a Gaussian distribution  $\mathcal{N}(35, \sigma^2 = 2)$ .
- ▶ Estimate the true temperature using MLE, MAP and MLE estimators

# Exercise

Exercise: constant value, 3 measurements, Gaussian same  $\sigma$

- ▶ What if he have three thermometers, showing 40, 38, 41 degrees

Exercise: constant value, 3 measurements, Gaussian different  $\sigma$

- ▶ What if the temperature this time of the year has Gaussian distribution  $\mathcal{N}(35, \sigma_2^2 = 3)$ 
  - ▶ different variance,  $\sigma_2 \neq \sigma$

# General signal in AWGN

- ▶ Consider that the true underlying signal is  $s_{\Theta}(t)$
- ▶ Consider AWGN noise  $\mathcal{N}(\mu = 0, \sigma^2)$ .
- ▶ As in Maximum Likelihood function, overall likelihood function

$$w(\mathbf{r}|\Theta) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{\sum (r_i - s_{\Theta}(t_i))^2}{2\sigma^2}}$$

- ▶ But now this function is also **multiplied with**  $w(\Theta)$

$$w(\mathbf{r}|\Theta) \cdot w(\Theta)$$

# General signal in AWGN

- ▶ MAP estimator is the argument that maximizes this product

$$\hat{\Theta}_{MAP} = \arg \max w(\mathbf{r}|\Theta)w(\Theta)$$

- ▶ Taking logarithm

$$\begin{aligned}\hat{\Theta}_{MAP} &= \arg \max \ln (w(\mathbf{r}|\Theta)) + \ln (w(\Theta)) \\ &= \arg \max -\frac{\sum (r_i - s_{\Theta}(t_i))^2}{2\sigma^2} + \ln (w(\Theta))\end{aligned}$$

# Gaussian prior

- ▶ If the prior distribution is also Gaussian  $\mathcal{N}(\mu_{\Theta}, \sigma_{\Theta}^2)$

$$\ln(w(\Theta)) = -\frac{\sum(\Theta - \mu_{\Theta})^2}{2\sigma_{\Theta}^2}$$

- ▶ MAP estimation becomes

$$\hat{\Theta}_{MAP} = \arg \min \frac{\sum(r_i - s_{\Theta}(t_i))^2}{2\sigma^2} + \frac{\sum(\Theta - \mu_{\Theta})^2}{2\sigma_{\Theta}^2}$$

- ▶ Can be rewritten as

$$\hat{\Theta}_{MAP} = \arg \min d(\mathbf{r}, s_{\Theta})^2 + \underbrace{\frac{\sigma^2}{\sigma_{\Theta}^2}}_{\lambda} \cdot d(\Theta, \mu_{\Theta})^2$$

# Interpretation

- ▶ MAP estimator with Gaussian noise and Gaussian prior

$$\hat{\Theta}_{MAP} = \arg \min d(\mathbf{r}, s_{\Theta})^2 + \underbrace{\frac{\sigma^2}{\sigma_{\Theta}^2}}_{\lambda} \cdot d(\Theta, \mu_{\Theta})^2$$

- ▶  $\hat{\Theta}_{MAP}$  is close to its expected value  $\mu_{\Theta}$  and it makes the true signal close to received data  $\mathbf{r}$ 
  - ▶ Example: “search for a house that is close to job and close to the Mall”
  - ▶  $\lambda$  controls the relative importance of the two terms
- ▶ Particular cases
  - ▶  $\sigma_{\Theta}$  very small = the prior is very specific (narrow) =  $\lambda$  large = second term very important =  $\hat{\Theta}_{MAP}$  close to  $\mu_{\Theta}$
  - ▶  $\sigma_{\Theta}$  very large = the prior is very unspecific =  $\lambda$  small = first term very important =  $\hat{\Theta}_{MAP}$  close to ML estimation

# Applications

- ▶ In general, practical applications:
  - ▶ can use various prior distributions
  - ▶ estimate **multiple parameters** ( a vector of parameters)
- ▶ Applications
  - ▶ denoising of signals
  - ▶ signal restoration
  - ▶ signal compression



# Estimator bias

- ▶ How good is an estimator?
  - ▶ Many ways to characterize
- ▶ An estimator  $\hat{\Theta}$  is a **random variable**
  - ▶ can have different values, because it is computed based on the received samples, which depend on noise
  - ▶ example: in lab, try on multiple computers => slightly different results
- ▶ As a random variable, it has:
  - ▶ an average value (expected value):  $E \{ \hat{\Theta} \}$
  - ▶ a variance:  $E \{ (\hat{\Theta} - \Theta)^2 \}$

- ▶ **Unbiased** estimator = if the average value of the estimator is the true value of  $\Theta$

$$E \{ \hat{\Theta} \} = \Theta$$

- ▶ **Biased** estimator = if the average value of the estimator is different from the true value  $\Theta$ 
  - ▶ the difference  $E \{ \hat{\Theta} \} - \Theta$  is called **the bias** of the estimator

# Estimator bias

- ▶ Example: for constant signal  $A$  with AWGN noise (zero-mean), ML estimator is  $\hat{A}_{ML} = \frac{1}{N} \sum_i r_i$
- ▶ Then:

$$\begin{aligned} E \{ \hat{A}_{ML} \} &= \frac{1}{N} E \left\{ \sum_i r_i \right\} \\ &= \frac{1}{N} \sum_{i=1}^N E \{ r_i \} \\ &= \frac{1}{N} \sum_{i=1}^N E \{ A + \text{noise} \} \\ &= \frac{1}{N} \sum_{i=1}^N A \\ &= A \end{aligned}$$

- ▶ This estimator is unbiased

# Estimator variance

- ▶ Unbiased estimators are good, but if the **variance** of the estimator is large, then estimated values can be far from the true value
- ▶ We prefer estimators with **small variance**, even if maybe slightly biased